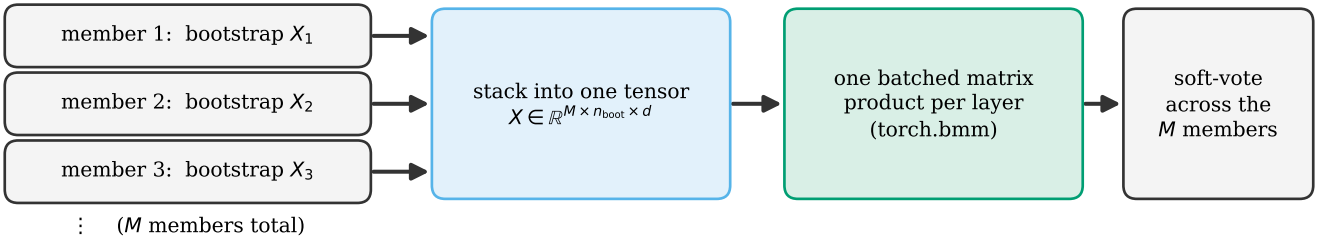
member 1: bootstrap X_1
member 2: bootstrap X_2
member 3: bootstrap X_3
⋮ (M members total)
stack into one tensor
$X \in \mathbb{R}^{M \times n_{\text{boot}} \times d}$
one batched matrix product per layer (torch.bmm)
soft-vote across the M members