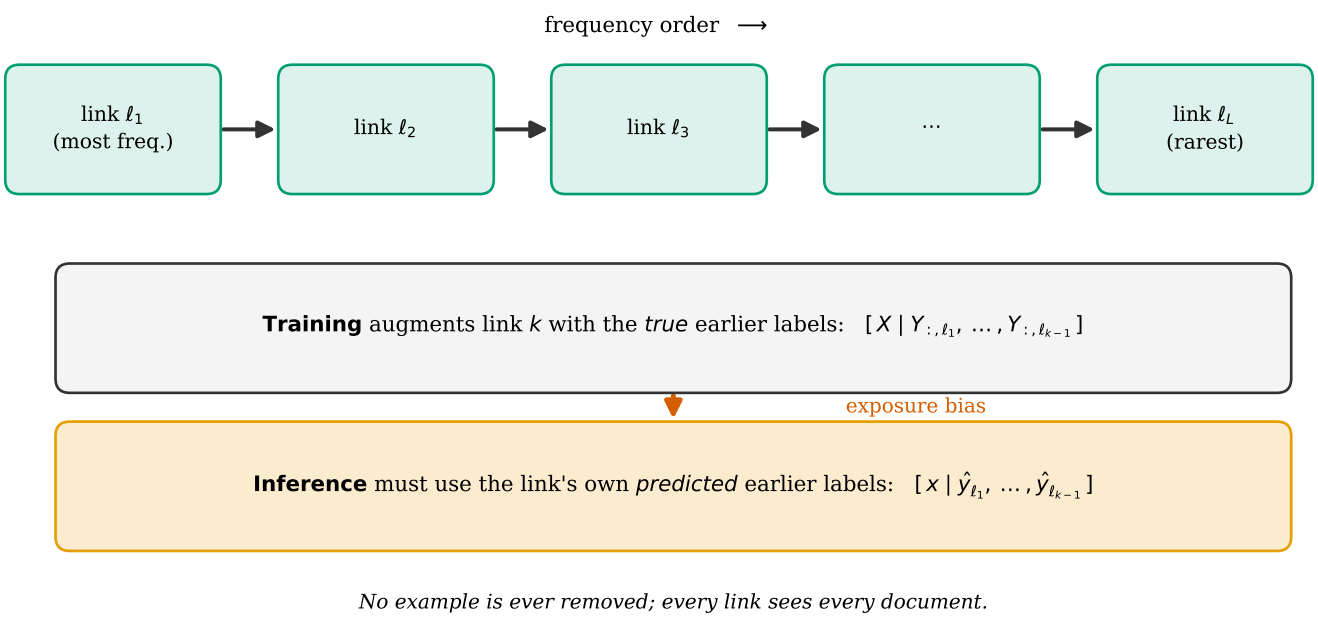
frequency order ⟶
link ℓ_1 (most freq.)
link ℓ_2
link ℓ_3
...
link ℓ_L (rarest)
Training augments link k with the *true* earlier labels: $[X \mid Y_{:,\ell_1}, \ldots, Y_{:,\ell_{k-1}}]$
exposure bias
Inference must use the link's own *predicted* earlier labels: $[x \mid \hat{y}_{\ell_1}, \ldots, \hat{y}_{\ell_{k-1}}]$
No example is ever removed; every link sees every document.